# Do Generative Large Language Models Close the Language Gap in Multilingual Discourse Classification? A Systematic Evaluation

Word count: 7.673

## Abstract

Uneven performance between high- and low-resource languages remains a central challenge for applying large language models to social-science text analysis. This paper evaluates whether generative models narrow that gap on two multilingual classification tasks – xenophobic and anti-government discourse – using 6,000 social media posts in Danish, Swedish, and German across 12 encoder, encoder-decoder, and decoder models under supervised and prompted regimes. The gap persists: model configurations applied to German (or English translations) outperform those for Danish and Swedish; fine-tuned encoders set the ceiling (macro-F1 ≈ 0.85 in German), while few-shot prompting yields only modest gains over zero-shot. A consensus analysis of per-item agreement shows higher model–gold alignment for xenophobia than anti-government (high-agreement band: 76.6% vs 52.3% of items), with residual errors concentrated in mediated discourse and irony/euphemism. Methodologically, I introduce a consensus-banding framework linked to a mixed-effects logistic model with crossed item/model intercepts, providing a reproducible bridge between quantitative reliability and qualitative error typology. Overall, language resources – more than model family or size – govern reliability in theory-driven multilingual discourse analysis, underscoring the need for evaluation practices and infrastructure that center low-resource languages.

# Introduction

Text classification is one of the most important tasks in natural language processing (NLP) and increasingly central to social science research (Vermeer et al. 2019). The task involves assigning text to predefined categories—topics, sentiments, discourse dimensions, emotions or frames—and underpins large-scale empirical analyses across platforms and domains. Until the 2022 release of ChatGPT most pipelines relied on human-labeled data to fine-tune encoder-based models such as BERT (Devlin et al., 2019). These models excel at learning contextual representations via masked language modeling but require costly annotation. The rise of large-scale generative models has introduced alternative approaches. Decoder-based models such as ChatGPT, Claude-Sonnet, Llama, or DeepSeek generate text by predicting next tokens autoregressively, enabling zero-/few-shot classification via prompting without fine-tuning (Kroon et al., 2024; Nielsen et al., 2025). A third family, encoder-decoder models like T5 (Raffel et al., 2020) or BART (Lewis et al., 2020), combine both architectures and can be either fine-tuned or prompted. While terms like “transformers”, “LLMs”, and “generative AI” are

sometimes used interchangeably, architectural distinctions matter for classification tasks. In this study, I evaluate all three model families—encoder-only, decoder-only, and encoder-decoder—to assess their performance and resource requirements for social science text classification.

Generative models promise to reduce annotation costs and have shown competitive performance on several classification benchmarks without task-specific training, including multilingual settings (e.g., Gilardi et al., 2023; Heseltine & Clemm Von Hohenberg, 2024; Törnberg, 2024). GPT-3.5 Turbo and GPT-4 Turbo outperform dictionary approaches and fine-tuned traditional machine learning classifiers on detecting sentiments, discrete emotions (Rathje et al., 2024), and political texts (Heseltine & Clemm Von Hohenberg, 2024) across multiple languages. Yet, three concerns remain salient. First, decoder-based models often must be substantially larger to match the accuracy of encoder-based models such as BERT on discriminative classification (Schick & Schütze, 2021). Second, models systematically perform better in high-resource languages due to richer pretraining and tooling, while low-resource languages lag behind (Robinson et al., 2023). According to Baden et al. (2022), multilingually comparative studies make up less than 16% of all top-ranking studies on communication, and 18% in political science compared to studies focusing on English. Third, the black-box nature of decoder-based models raises concerns about measurement validity and reproducibility (Baden et al., 2022). And even when outputs appear valid, we still lack clarity on how LLMs arrive at decisions and which biases they may introduce during classification (Laurer et al., 2025).

This paper addresses all three issues with a systematic evaluation of multilingual, political discourse-level classification (Grimmer & Stewart, 2013). I study two substantively important constructs of political discourse – xenophobic and anti-government discourse – both

of which can appear in overt and implicit forms, and where reliable measurement requires sensitivity to pragmatic cues, coded language, and shifting rhetorical strategies (Cisneros & Nakayama, 2015; Matamoros-Fernández & Farkas, 2021). As a result, detecting such discourse challenges both large language models and social scientists, who must account for cultural variations, subtle bigotry, and evolving political narratives.

These questions are central for communication research. Cross-national studies increasingly rely on automated text classification to track the circulation of far-right narratives, anti-institutional sentiment, and hostile speech across platforms and languages. If model performance varies systematically by language and discourse type, this can bias estimates of issue salience, misrepresent the prevalence of xenophobia or anti-government rhetoric, and undermine the comparability of findings across media systems. Multilingual evaluation of LLM-based classifiers is therefore not only a technical matter, but a question of measurement validity and cross-lingual comparability in empirical communication research.

Relative to prior multilingual LLM evaluations, this study contributes in three ways. First, it offers an architecture-aware comparison that crosses encoder-only, decoder-only, and encoder–decoder models with both supervised fine-tuning and zero-/few-shot prompting, including a small, supervised decoder condition to reduce the architecture–inference confound. Second, it centers low-resource Scandinavian languages (Danish, Swedish) and discourse-level constructs (xenophobia, anti-government rhetoric) drawn from right-wing social media, extending multilingual evaluation beyond English and high-resource benchmarks. Third, it combines model-level metrics with a consensus-based error audit: treating model configurations as independent raters, banding items by agreement, and coding low-consensus cases into a discourse-informed error typology. This mixed-methods design links quantitative

performance differences to specific linguistic and pragmatic failure modes. This leads to two main research questions:

> **RQ1**: How do different model families (decoder-based, encoder-based, and encoder-decoder) perform on political discourse classification across languages with varying resource levels?
>
> **RQ2**: Which linguistic features and discourse characteristics (e.g., irony, euphemism, mediated discourse, and role ambiguity) drive disagreement between models and expert coders?

More specifically, I focus on Danish and Swedish (low-resource languages), German (a high-resource language) and English translations to test whether decoder-based models close the gap relative to supervised encoders. I use *high-resource* and *low-resource* in their technical NLP sense: they refer to the relative availability of digital text corpora, pretrained models, and labeled datasets and tools for a given language. English and German are typically high-resource in this sense, whereas Danish and Swedish currently have substantially fewer training corpora, benchmarks, and off-the-shelf models and are thus treated as low-resource. Across 6,000 social media posts sampled from right-wing affiliated accounts (January 2019 – June 2022), I evaluate 12 models spanning encoder, encoder–decoder, and decoder families under zero-shot, few-shot, and supervised regimes. The study combines standard performance metrics (macro-F1 by language, task, and architecture) with a consensus-based error analysis that treats model configurations as independent raters, bands items by agreement, and qualitatively codes low-agreement posts into a six-category error taxonomy (mediated discourse, actor-role conflation, numbers-as-argument, irony/sarcasm, euphemism/dog-whistling, other/ambiguous).

Three results stand out. First, the language gap persists: German and English-translation runs outperform Danish and Swedish across families and prompting strategies. Second, smaller, supervised encoders (e.g., mDeBERTa/mmBERT) outperform much larger generative models on classification of both anti-government and xenophobia, while few-shot prompting delivers only modest gains over zero-shot. Third, error audits reveal that mediated discourse (reported speech/quotation) and irony/sarcasm drive much of the residual disagreement, highlighting where discourse-level inference remains brittle. The study contributes with (i) a controlled, architecture-aware comparison that separates prompting gains from supervision; (ii) multilingual evidence centered on low-resource languages and discourse-level constructs; and (iii) a consensus-guided qualitative analysis that pinpoints recurrent sources of error. Substantively, the results indicate that, in this setting, language resources appear to matter at least as much as – and often more than – model family or size for reliable, theory-driven multilingual classification, and they underscore the need for evaluation practices and infrastructure that prioritize low-resource settings while remaining attentive to task- and context-specific variation.

## Literature review: Encoders vs. decoders for multilingual, discourse-level classification

LLMs have opened at least two distinct pipelines for text classification in the social sciences by enabling both supervised fine-tuning of encoder-based models (e.g., BERT) and prompt-based inference with generative decoders (Gilardi et al., 2023; Kristensen-McLachlan et al., 2025; Zhao et al., 2023). Yet, while these models have proven capable in text classification, uncertainties remain regarding under what design choices they accurately and reliably do so – especially in multilingual and discourse-level settings where constructs are subtle and context dependent (Nielsen et al., 2025).

### *Encoders, Decoders and Encoder-decoders*

Encoder-only models (e.g., BERT, mBERT, XLM-R) are pretrained with masked-language modeling to learn bidirectional contextual representations. When fine-tuned on annotated data, they adapt a classification head to the target construct and provide high precision/recall on discriminative tasks (Conneau et al., 2020; Devlin et al., 2019). This characteristic enables them to achieve high precision when training and inference data originate from the same or highly similar data sources yet potentially limits their adaptability when faced with new or unseen linguistic contexts.

Decoder-only models (GPT-class) are pretrained autoregressively and often instruction-tuned (e.g., reinforcement learning from human feedback, enabling zero-/few-shot classification via prompting without parameter updates (Brown et al., 2020; Ouyang et al., 2022). Decoder-based models hold two advantages: First, their often-large scale impart broad world knowledge and stylistic flexibility, enabling them to generate human-like text across domains. Second, they exhibit in-context learning: by framing a handful of examples or explicit instructions within the prompt, they can "generalize" to new tasks—such as detecting xenophobic or anti-government discourse—without parameter updates (Dong et al., 2024; Gu et al., 2023).[1] However, for discriminative classification they generally require substantially larger sizes to match BERT-family accuracy, and decision boundaries induced purely by prompts can underspecify edge cases; multiple comparative studies find that fine-tuned

[1] The first decoder-based model—the Generative Pre-trained Transformer (GPT-1)—was developed by Radford and authors (2018) to accomplish downstream tasks with fine-tuning, but also—somewhat to the surprise of the developers—performed well on zero-shot performance on question-answering, schema resolution, sentiment analysis (Nassiri & Akhloufi, 2023).

encoders still hold an accuracy edge on precise classification (Bucher & Martini, 2024; Nielsen et al., 2025; Schick & Schütze, 2021). Encoder-decoder (seq2seq) models (e.g., T5/mT5) sit between these paradigms: an encoder builds a representation, a decoder generates a label string. For pure classification they can behave like encoders under supervision but are often less parameter-efficient in this setting (Xue et al., 2021).

### *Discourse Classification with LLMs in Multilingual Contexts*

Discourse classification requires reasoning beyond sentence boundaries—tracking reference, rhetorical relations, dialog state, and implicit content—capabilities that standard token-level signals only weakly capture (Y. Zhu et al., 2024). Core NLP subcomponents such as coreference and discourse parsing remain uneven across languages and domains, which constrains downstream classification that depends on them.

Multilingual encoders (e.g., mBERT or XLM-R) have demonstrated progress in zero-shot classification scenarios where models trained on pivot languages generalize effectively to unseen languages (Pires et al., 2019; Wu & Dredze, 2020). However, performance across languages is notably uneven, particularly for low-resource languages or those typologically distant from training languages (Lauscher et al., 2020; Wu & Dredze, 2020). In practice, even strong multilingual models underperform on under-represented languages, reinforcing the very "language-resource gap" this paper targets (Choenni et al., 2023). Furthermore, very little research has been carried out on discourse understanding and classification with LLMs in multilingual contexts (Eichin et al., 2025) and evaluation of multilingual LLMs (Choenni et al., 2024). Previous research underscores substantial gaps in empirical evaluations of multilingual classification for complex classification tasks especially in languages other than English (Nielsen et al., 2025).

Using both mono- and multilingual models, Mets and colleagues (2024) have conducted a comparative evaluation of supervised and zero-shot stance classification models on Estonian (low-resource language) immigration discourse. They find that even the best-performing classifiers struggled with key discourse-level challenges, including sarcasm, pragmatically ambiguous or context-dependent statements, and indirect or mediated speech. These cases produce systematic misclassification, particularly for implicit hostility (dog-whistles) and for political criticism whose target is context-dependent (L. Wei et al., 2025). This aligns with broader work on implicit hate speech, which repeatedly shows that coded or veiled xenophobic content is substantially harder to detect than explicit slurs or overt hostility (ElSherief et al., 2021; Kim & Lee, 2025; Ocampo et al., 2023).

A common mitigation in multilingual pipelines is source-side machine translation to English, then applying stronger English models (Windsor et al., 2019). Evidence suggests this can be practically effective and reproducible when paired with encoders (e.g., BERT), but benefits are mixed and task-dependent (Artetxe et al., 2023; Licht, 2023); machine translation adds its own biases and reproducibility trade-offs (API cost/opacity). Translation can raise a floor but rarely closes architecture- or resource-driven gaps by itself – consistent with recent small, mixed gains (Licht et al., 2024). Finally, early cross-lingual comparisons between compact multilingual encoders and prompted decoders show that, for precise discriminative classification, fine-tuned encoders often retain an edge, while decoders offer adaptability with modest few-shot gains (X. Zhu et al., 2024).

Taken together, the literature emphasizes (i) persistent gaps by language resource level (low- vs. high-resource languages), (ii) that implicit, coded, and context-dependent forms of both hostility and political stance are systematically harder to classify than more explicit, and (iii) limited but sometimes useful gains from translation and few-shot prompting.

### *Prompting Strategies for Zero- and Few-Shot Classification*

Prompting governs how decoder LLMs instantiate a decision rule at inference time, and small design choices (definitions, exemplars, decoding) can shift both accuracy and failure modes (Weber & Reichardt, 2023). Recent studies converge on three practical pillars: specify the construct, constrain the output, and validate against gold data. First, prompts should explicitly define target categories in domain language, include inclusion/exclusion rules, and—when using few-shot—provide stratified, label-balanced exemplars that match the deployment languages and genres (Farjam et al., 2025; Weber & Reichardt, 2023). This reduces label drift – the tendency for category boundaries to shift or be applied inconsistently across cases – and improves reproducibility in annotation pipelines.

Second, structure the interaction to minimize format errors and spurious variation: determinize decoding (e.g., low/zero temperature) (Renze & Guven, 2024), cap generation length, and force a schema (e.g., a single JSON field for the label plus a short rationale). Empirically, temperature often has weak or task-specific effects on accuracy, so defaulting to near-deterministic decoding is a safe baseline for measurement settings (Guo et al., 2023).

Third, select prompting strategies knowing both gains and risks. Regarding zero vs. few-shot prompting, existing research suggests that few-shot classification, which incorporates minimal labelled examples, tends to outperform zero-shot methods marginally, particularly for complex social science tasks like political and sentiment analyses (Kristensen-McLachlan et al., 2025). Although chain-of-thought (CoT) prompting can improve multi-step arithmetic/logic benchmarks and zero-shot reasoning in large decoders, evidence is mixed for short, binary text classification; CoT can also add variance and its explanations are not necessarily faithful to the model's decision process (J. Wei et al., 2023). When using rationales,

keep them short and non-speculative, and prefer "explain after labeling" (for auditing) over "explain to decide" unless you can sandbox safety impacts (Kojima et al., 2022).

For multilingual discourse classification, at least two additional choices matter. *First*, prompts should be construct-invariant across languages: keep the same label inventory and as-close-as-possible translations of category definitions and instructions so that "xenophobic" or "anti-government" discourse denote the same underlying construct in Danish, Swedish and German. Within that shared template, example selection should be language-aware: instantiate few-shot examples in the same language as the input (Danish exemplars for Danish texts, etc.), rather than relying only on English exemplars and hoping for cross-lingual transfer. This helps models pick up language-specific cues (e.g., irony markers, dog-whistles) without changing the underlying definition. *Second*, if translation to English is used for convenience, treat it as a separate condition: apply the same prompt template to the translated inputs, but do not mix exemplars or outputs across languages during evaluation. Recent guidance for social-science coding stresses reporting prompt text, exemplar provenance, decoding parameters, and retry logic to support reproducibility (Weber & Reichardt, 2023).

Taken together, the prompting literature suggests a pragmatic baseline for decoder-based classification pipelines: start with definition-first zero-shot prompts plus explicit output constraints; add few-shot exemplars only when they are in-domain and language-matched; keep decoding deterministic (e.g., low temperature) and enforce a simple schema (such as a single JSON label); report prompt text and parameters for reproducibility; and validate against gold annotations head-to-head with supervised baselines.

## Materials and Methods

### *Sampling – right-wing social media accounts in German, Danish and Swedish*

The study draws on a stratified sample of 32,983 posts in Danish, German, and Swedish from a larger corpus of 42 million posts by ~31,000 right-wing accounts (January 2019 - June 2022). Accounts were identified via mutual sharing within a right-leaning information-sharing network (see AUTHORS, 2024 for a similar approach). Platforms span Twitter/X, Facebook, Telegram, YouTube, Gab, Reddit, and VKontakte. All 32,983 posts were manually annotated for xenophobic and anti-government discourse. For model comparison, I reserved a held-out evaluation set of 2,000 posts per language (n = 6,000), while the remaining 26,983 annotated posts were used exclusively for supervised training and validation.[2]

### *Annotating xenophobic and anti-government discourse*

Four expert coders annotated the data for xenophobic and anti-government discourse (see Appendix A1 for codebook). In this context, *xenophobia* is understood as fear, hatred or hostility regarding perceived foreigners (Mudde, 1995) and *anti-government rhetoric* is understood as outspoken distrust of governmental institutions and policies (Campion et al., 2023). The coding was binary (present or not present) to ensure consistency and simplicity.

The two discourse dimensions were selected to span a range of task difficulty in the corpus. In the data, xenophobic discourse more often includes overt lexical cues (e.g., slurs, exclusionary formulas), which can be expected to be comparatively easier for LLMs to detect. At the same time, xenophobia also appears in more implicit, euphemistic, or dog-whistled forms. By contrast, much of the anti-government rhetoric in this sample is expressed through context-dependent references to country-specific institutions, policies, and actors: identical

[2] See ANONYMIZED for annotated dataset.

surface forms can signal satire, policy critique, or generalized distrust depending on national context. These patterns led me to the expectation that anti-government labels would be more sensitive to contextual cues than xenophobic ones, even though both constructs can be expressed overtly or implicitly. This raises the risk that models import domain knowledge from other settings and misclassify locally specific references. This design choice ensures the evaluation spans a meaningful difficulty gradient rather than a single construct, allowing for probing where multilingual reliability breaks down. In the 6,000-post evaluation dataset, xenophobic content was present in 9.32% (n=559) of the posts while anti-government rhetoric appeared in 9.28% (n=557) of the posts, indicating similar distributions of these narratives across the sample.

Intercoder reliability was assessed on a stratified sample of 1,000 posts (balanced by language and discourse dimension), coded independently by all four annotators. To facilitate cross-language comparison for reliability only, items were machine-translated to English; all modeling used native-language texts. Krippendorff's α (nominal) was 0.77 for xenophobia and 0.71 for anti-government (Krippendorff, 2004). Disagreements were adjudicated via a pre-specified protocol to produce a single gold label per item (see Appendix A2 for tests).

Minimal preprocessing was applied; links were removed and emojis were converted to text, while stop words, hashtags, and mentions were retained. As machine translation into English is often recommended to leverage the availability of larger and higher-quality English-trained models (Laurer et al., 2023; Máté et al., 2023), posts were translated to benchmark against the multilingual models. Translations were carried out using DeepL API.

Two potential threats to external validity merit attention. First, platform-specific language may affect classification accuracy. Platforms that allow for longer or more

multimodal expressions (e.g., Gab and Telegram) often contain in-group memes, coded slang, or other stylistic markers that differ from the shorter posts or comments on X/Twitter. Classifier performance on such registers may therefore be under- or over-estimated. Second, temporal drift poses a challenge: the sampling period overlaps with the COVID-19 pandemic and the emergence of new right-wing slang (e.g., plandemic, globalists), raising questions about stability of model performance over time.

### *Models overview*

The evaluated mix of encoder-, decoder-based and encoder-decoder LLMs was selected based on performance from EuroEval's natural language understanding (NLU) leaderboard (see Table 1 for models). The models were selected to attain maximum variance regarding past classification performance on a range of tasks, accessibility, usage in research and price.

**Table 1.** Overview of language models

| Model name | Architecture | Access | Classification | Note |
|---|---|---|---|---|
| **Gemma-3-27b (Google DeepMind)** | Decoder-based | Open-source | Zero/Few-shot with Huggingface | |
| **Llama-70b-4bit (Meta)** | Decoder-based | Open-source | Zero/Few-shot with Huggingface | |
| **Llama-3.2-3b (Meta)** | Decoder-based | Open-source | Fine-tuning with Huggingface | |
| **Gemma2-2b (Google DeepMind)** | Decoder-based | Open-source | Fine-tuning with Huggingface | |
| **mDeBERTa-276m (Microsoft)** | Encoder-based | Open-source | Fine-tuning with Huggingface | Uses NLI reformulation |
| **Multilingual Modern BERT-110M (mmBERT)** | Encoder-based | Open-source | Fine-tuning with Huggingface | Uses "direct sequence classification" |
| **mT5-large (Google)** | Encoder-decoder | Open-source | Fine-tuning with Huggingface | Text-to-text (seq2seq) NLI reformulation for classification |
| **Gemini-2-flash-lite (Google)** | Decoder-based | Proprietary | Zero/Few-shot with Openrouter | $0.075/M input tokens $0.30/M output tokens |

| | | | | |
|---|---|---|---|---|
| **Gemini-2-flash (Google)** | Decoder-based | Proprietary | Zero/Few-shot with Openrouter | $0.10/M input tokens $0.40/M output tokens |
| **GPT-4o (OpenAI)** | Decoder-based | Proprietary | Zero/Few-shot with Openrouter | $3.75/M input tokens $15/M output tokens |
| **GPT-4o-mini (OpenAI)** | Decoder-based | Proprietary | Zero/Few-shot with Openrouter | $0.30/M input tokens $1.20/M output tokens |
| **Claude Sonnet-3.5 (Anthropic)** | Decoder-based | Proprietary | Zero/Few-shot with Openrouter | $3.0/M input tokens $15/1M output tokens |

**Notes.** Unless noted otherwise, all random seeds = 42; fine-tuned models were trained/evaluated on a single NVIDIA A100 (80 GB). Proprietary decoders were accessed via OpenRouter; open-source decoders ran locally.

## *Training and prompting strategies*

### Supervised fine-tuning (encoder and encoder–decoder)

The mDeBERTa-v3 models used for supervised text classification were also prompted using a natural language inference (NLI) approach using "hypothesis"-pairs for optimized performance (Laurer et al., 2022) (see Table 2). NLI models are trained to decide whether a short statement ("hypothesis") is supported by a given text ("premise"), so I reformulate each of the labels as such hypotheses (e.g., "This post expresses xenophobic discourse") and train the model to judge whether they hold for each post. For fine-tuning of the mDeBERTa-v3, T5 and mmBERT models, the annotated dataset without the evaluation dataset to avoid data leakage (N = 26,983 posts) were used. Splits were 80/20 (train/test; `random_state=42`) with the model parameters updated through gradient descent optimization across multiple epochs adopting standard hyperparameters to ensure comparability and reproducibility (Sun et al., 2020). To mirror class imbalance while stabilizing learning, I under-sampled negatives in the *training* portion only so that positives ≈ 20%; evaluation distributions remained intact (Hoyos-Osorio et al., 2021). mDeBERTa received paired inputs *(post, hypothesis)*; mmBERT used single-sequence inputs; mT5 used seq2seq prompts with short, literal targets. Hyperparameters, batching, and save/eval schedules follow the model-specific settings in Appendix B1.

### PEFT fine-tuning of decoder-based models

In addition to large decoder-only LLMs used purely via prompting, I also include a supervised decoder condition to reduce the architecture–inference confound. Two small decoder-only models (Gemma2-2B and Llama-3.2-3B) were therefore fine-tuned on the same 80/20 splits and class balance as the encoder models, using LoRA adapters ($r = 32$, $\alpha = 64$, dropout = 0.10) with frozen base weights. These models remain members of the decoder family; the difference is not architectural but the inference regime: instead of using prompts to induce a decision rule at run time, they learn a classification head under supervision, analogous to the encoder baselines. This allows me to ask whether the accuracy advantage of supervised models over prompted decoders is specific to encoder architectures or also appears for decoder-only models of comparable size.

Optimization used AdamW (weight decay = 0.01), learning rate = 1e-4, 5 epochs, warm-up ratio = 0.10, cosine LR schedule, gradient checkpointing, max grad-norm = 0.5, and bf16 when available (else fp16). Per-device batch size was 8 with gradient accumulation = 2 (effective batch = 16). To address imbalance, I computed class weights from the train split and applied a focal loss ($\gamma = 2$) with a +50% boost to the positive-class weight; no label smoothing was used. Sequence length was 512 tokens with head/tail-preserving truncation. Evaluation ran every 100 steps with load-best-model-at-end (metric: F1 for the positive class). For these fine-tuned decoders, inference produced discrete labels via the classification head; I did not additionally prompt them, because the study already includes decoder-only models in a prompt-only condition, and combining prompting with a trained classification head would blur the contrast between supervised and prompt-based decoding that RQ1 aims to evaluate (full details in Appendix B2).

### *Prompted decoding (proprietary and large open-source decoders)*

The decoder-based models were prompted following best practice in prompt engineering describing key concepts, categories, coding rules, and preferably coded examples (Sivarajkumar et al., 2024). Table 2 gives the task definitions used in zero-shot and the few-shot structure for xenophobic discourse (anti-government omitted and few-shot prompt truncated for space; full templates in Appendix C). The mDeBERTa NLI prompts appear in the table for completeness but note that encoders were trained under supervision; prompted decoding was only applied to the decoder-only models listed above.

**Table 2.** Prompting strategy for xenophobic discourse

| **Prompt strategy** | **Xenophobic discourse** |
|---|---|
| Zero-shot classification | "**system role**": "You are a multilingual text classifier. Follow the task definition precisely. Do not explain, hedge, or refuse. Output must be valid JSON only.",.<br><br>"**content**": "Classify whether the post expresses xenophobic discourse (fear/hostility toward foreigners, immigrants, or out-groups) via stereotypes, dehumanization, exclusionary policy advocacy, or generalized negative attributions to groups defined by nationality/ethnicity.<br>Do not label as xenophobic: neutral news headlines, quotations/reporting of others' views without endorsement, criticism of specific policies without attributing negative traits to an out-group, or discussion of religion/culture without derogation.<br><br>Ambiguity rule<br>If evidence is insufficient or only implied, return 0 (not present).<br><br>Edge cases<br>- Mediated/quoted speech ("X says …"): require clear endorsement; otherwise 0.<br>- Irony/sarcasm/rhetorical questions: label 1 only if the target and derogation are unambiguous.<br>- Hashtags/links/usernames/emojis: may inform stance; do not use author metadata (handles, follower counts)."<br><br>"**role**": "user",<br>"**content**": f" "Is this text xenophobic? Reply with ONLY valid JSON: {{\"label\": true}} for present, and {{\"label\": false}} for not present. Do not explain.\n\nText: {text}", |

| | |
|---|---|
| Few-shot examples (same prompt as for zero-shot classification) | Same system role and content prompt as above + three examples:<br><br>“**role**”: “**user**”,<br>“**content**”: “Example 1: Asylum seeker shows his "gratitude." Muhammed assaulted 9-year-old Kiki, so maybe there are some who think that kind of thing is okay?”<br>},<br>{“**role**”: “assistant”, “**content**”: “true”},<br>{<br>“**role**”: “**user**”,<br>“**content**”: “Example 2: Southern-looking man: I had also been curious once about what is meant by a southern-looking man, as frequently mentioned in police reports. In effect, such offender descriptions significantly reduce the success of investigations. A scoundrel who thinks ill of it.”<br>},<br>{“**role**”: “assistant”, “**content**”: “true”},<br>{<br>“**role**”: “user”,<br>“**content**”: “Example 3: These immigrants are destroying our society and our culture. They should be sent back.”<br>},<br>{“**role**”: “assistant”, “**content**”: “true”}<br>{<br>“**role**”: “user”,<br>“**content**”: f”Is this text xenophobic? Reply with ONLY valid JSON: {{\"label\": true}} for present, and {{\"label\": false}} for not present. Do not explain.\n\nText: {text}”, |
| Encoder fine-tuning (NLI) | *Xenophobia: No:* "This quote is neutral or positive towards foreigners, immigrations, strangers or people from different cultures"<br><br>*Xenophobia: Yes:* "This quote is negative towards foreigners, immigrations, strangers or people from different cultures" |

**Note:** See Appendix C for full prompt overview. Decoder-based models (Gemma2 and Llama 3.2) are typically not prompted when fine-tuned for classification tasks.

## *Evaluation metrics for the quantitative analysis*

For RQ1, I evaluate each classifier against gold labels for binary xenophobia and anti-government detection. Macro-F1 was the primary evaluation metric due to class imbalance (~10% positive). I compute macro-F1 as the unweighted mean of the positive- and

negative-class F1 scores using `scikit-learn` and also report micro-F1 (average="micro") and accuracy.

For error profiling I include class-specific precision/recall/F1. Ground-truth labels and model outputs are standardized to binary {0,1}, rows with missing predictions are excluded, and per-language metrics are computed only when ≥10 valid items are available. Encoder and fine-tuned decoder models use argmax over a binary head with a 0.5 threshold. Generative decoders are constrained to true/false with temperature 0.0; one retry is issued on format errors, and if still invalid the response is counted as incorrect. Few-shot exemplars are drawn only from the training partition.

In Appendix D, I report detailed metrics overall for both xenophobic and anti-government discourse; English rows reflect DeepL translations. For summary tables by model category/strategy/resource level, the mean macro-F1 is reported with 95% confidence intervals computed as mean ± $t(n-1, 0.975) \times SE$.

### *Evaluation metrics for qualitative analysis: Consensus-based validation and error analysis*

For RQ2, I treat each model configuration as an independent "rater" and the human label as gold (see Appendix E for details). This follows a mixed-methods convergent validation logic (Creswell & Plano, 2022): quantitative agreement across models guides case selection, which is then examined qualitatively to identify recurring failure types.

For each test item, I assemble a wide prediction table with one column per classifier (zero-/few-shot decoders, supervised encoders/decoders/encoder–decoders) and parse all outputs to strict binary labels. Malformed or missing outputs are dropped for that classifier–item pair, and I report overall and language-specific missingness. I compute an agreement ratio per item as the share of *valid* model predictions matching the gold label:

$$\text{Agree}(i) = \frac{1}{M_i} \sum_{m=1}^{M_i} 1(\ \hat{y}_{im} = y_i$$

where $M_i$ is the number of valid predictions for item $i$.

Item-level agreement is defined as the proportion of valid model predictions that match the gold label. Items with fewer than three valid predictions are excluded. For descriptive diagnostics, items are grouped into three agreement bands: low (< .70), mid (.70–< .90), and high (≥ .90), following best practices in multi-rater assessment and outlier detection (Rousseeuw & Leroy, 1987). I summarize agreement by task (xenophobia vs. anti-government) and language (Danish, German, Swedish) and also report language-stratified band shares. To characterize variability without over-interpreting single models, I report mean macro-F1 and t-based 95% confidence intervals over configurations within each (architecture × evaluation method × language × task) cell. In addition, I fit a mixed-effects logistic regression (DV: model matches gold) with fixed effects for task and language and crossed random intercepts for items and models to test whether agreement differs systematically by task and language.

To probe systematic failure modes, I focus on low-agreement items (agreement < .70) and conduct a content analysis on this subset (n = 1,727). Two independent coders apply a pre-specified taxonomy with six error types:

1. **Mediated discourse/reporting** – quotation or reported speech (e.g., "X says that Y did Z").
2. **Actor–role conflation** – confusing non-government actors with government actors, especially in anti-government discourse.
3. **Numbers-as-argument** – implicit stance expressed through statistics or counts, common in xenophobic discourse.
4. **Irony/sarcasm** – satirical, mocking, or ironic tone.
5. **Euphemism/dog-whistling** – implicit or coded language.

6. **Other/ambiguous** – unclear or miscellaneous cases.

The taxonomy was developed iteratively: mediated discourse, irony/sarcasm, and euphemism/dog-whistling reflect commonly reported failure modes in text classification, while actor–role conflation and numbers-as-argument were added after initial coding revealed their prevalence in this dataset. Intercoder reliability for error-type coding was substantial ($\kappa$ = 0.737; 73.7% agreement), with highest agreement for xenophobic discourse in Danish (80.1%) and lowest for anti-government discourse in Danish (68.8%). Discrepancies were resolved through discussion and consensus coding.

This layered design serves three purposes: (i) to flag potential inconsistencies in the annotated gold labels via unusually low or high consensus; (ii) to isolate linguistic and contextual features that systematically challenge automated classification; and (iii) to assess cross-lingual performance variability through language-stratified agreement patterns. Posts at the extremes of the agreement distribution are then examined qualitatively to deepen interpretation of classifier behavior and sources of systematic error in political discourse classification.

# Results

## *Quantitative assessment: Architectural differences, languages and prompt strategies*

The first part of the analysis investigates how choices of architecture—encoder-only, decoder-only, and encoder-decoder—impacts performance between different model types across low-, high-resource and translated languages (RQ1). Table 3 compares architectures by language resource level (Appendix D: Table D1 statistical tests; Table D2 model overview).

Performance varies by language resource and architecture. In high-resource settings (German and English translations), encoder models achieve notably higher macro-F1 (0.85

[0.70–0.96]) than decoder families: proprietary decoders are moderate (≈0.69–0.70) and open-source decoders lower (≈0.63–0.64). Across 33 configurations, the architectural advantage is most pronounced in German, where encoders consistently outperform decoders by substantial margins. Translation to English yields modest gains in some zero-/few-shot decoder runs but does not improve fine-tuned decoders. Proprietary decoders generally outperform open-source decoders across conditions, with the largest gaps in German (≈0.69–0.70 vs ≈0.63–0.64). Within architectures, few-shot vs. zero-shot differences for decoders are small and do not approach supervised encoder performance on the anti-government task.

**Table 3.** F1-performance on political discourse classification

| Model Type | Strategy | Low-resource | High-resource | Translated | Overall |
|---|---|---|---|---|---|
| **Decoder (PRO)** | Zero-shot | 0.65 [0.60-0.71] | 0.69 [0.64-0.73] | 0.70 [0.66-0.73] | 0.68 [0.64-0.72] |
| **Decoder (PRO)** | Few-shot | 0.65 [0.61-0.69] | 0.70 [0.65-0.75] | 0.69 [0.64-0.74] | 0.68 [0.63-0.72] |
| **Decoder (OS)** | Zero-shot | 0.59 [0.25-0.87] | 0.63 [0.26-0.98] | 0.65 [0.34-0.89] | 0.63 [0.30-0.90] |
| **Decoder (OS)** | Few-shot | 0.61 [0.48-0.73] | 0.64 [0.55-0.74] | 0.67 [0.49-0.83] | 0.64 [0.50-0.78] |
| **Decoder (OS)** | Supervised | 0.68 [0.58-0.78] | 0.70 [0.63-0.78] | 0.61 [0.23-0.91] | 0.65 [0.41-0.85] |
| **Encoder (OS)** | Supervised | 0.79 [0.63-0.92] | 0.85 [0.70-0.96] | 0.79 [0.72-0.86] | 0.80 [0.69-0.90] |
| **Encoder-decoder* (OS)** | Supervised | 0.72 | 0.76 | 0.70 | 0.72 |

**Note:** Macro-F1 (mean [95% CI]) combining xenophobia and anti-government. OS: open-source; PRO: proprietary. Low-resource: Danish/Swedish; High-resource: German; Translated: English. Wide CIs reflect small architecture Ns (n=1–2). *CI omitted (single model).

Encoders' advantage aligns with their bidirectional context modeling and supervised specialization. Fine-tuned encoders (mDeBERTa/mmBERT) learn task-specific patterns (e.g., xenophobic dog-whistles, anti-government markers) that prompting alone does not reliably capture. Encoder–decoders (T5/mT5) form a variable middle (overall ≈0.72), below encoders despite identical supervision.

These aggregate findings suggest architectural hierarchies; however, small per-architecture Ns warrant cautious interpretation. Figure 1 summarizes language–task

patterns (grid); Appendix D3 provides class-1 F1, which further strengthens the encoder advantage—especially on anti-government in German.

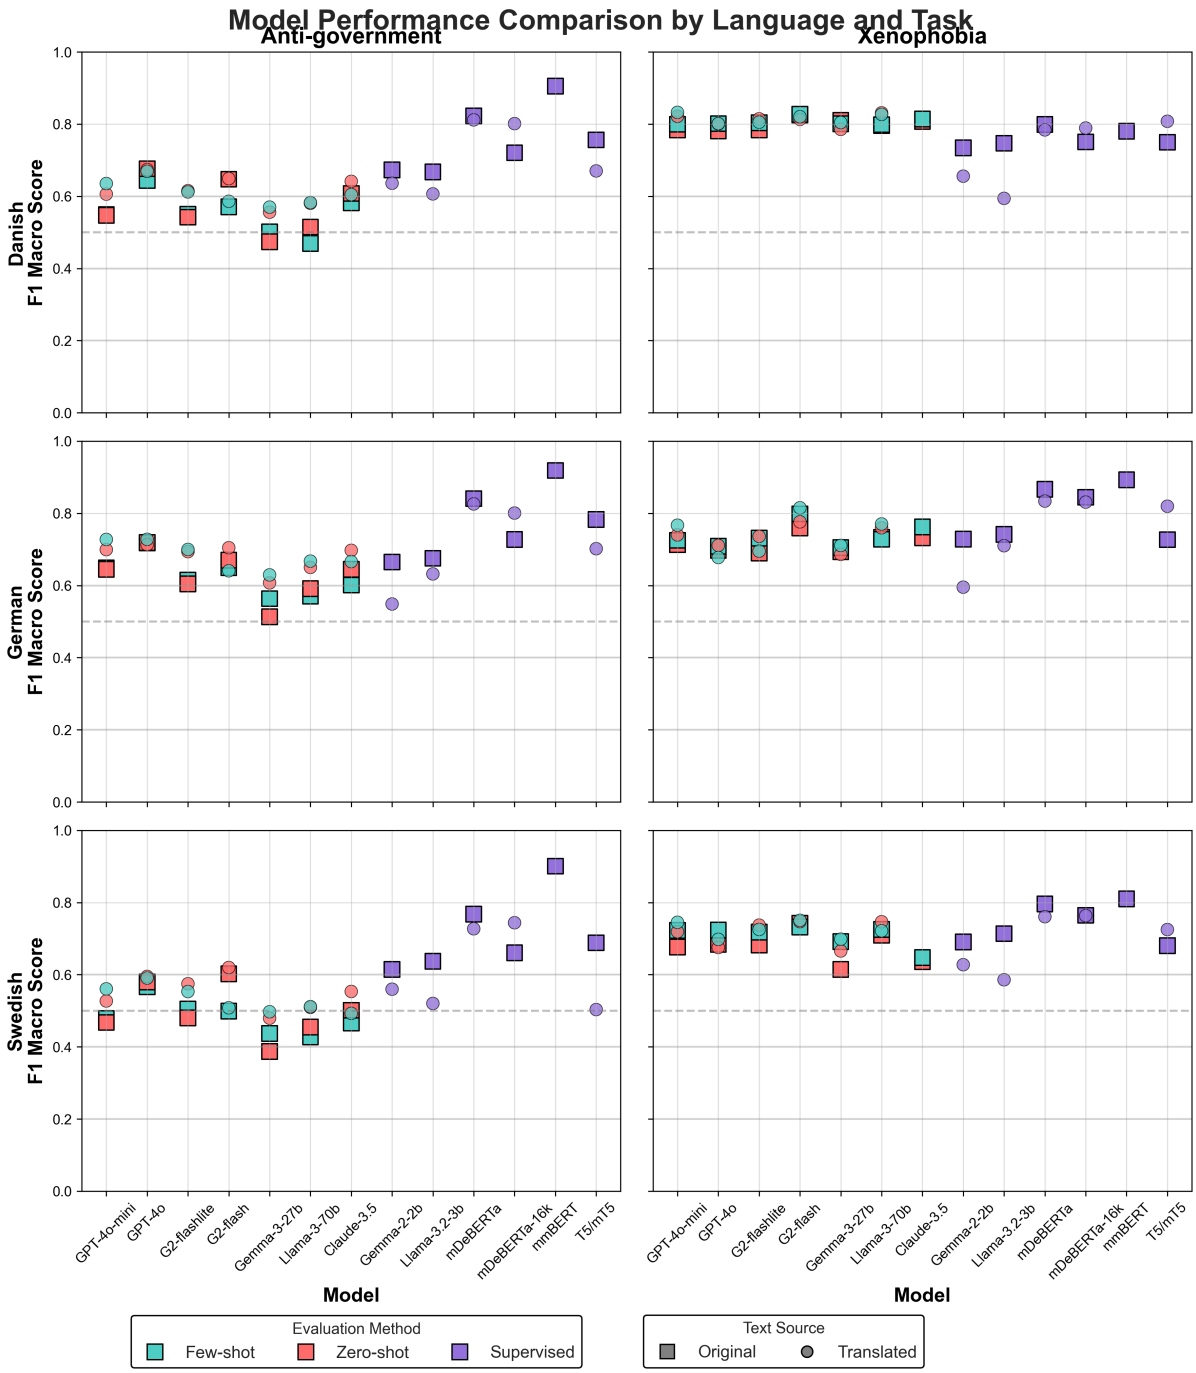


**Figure 1.** Cross-Linguistic Performance of Model Architectures

**Note:** Macro-F1 by architecture and evaluation method (rows: languages; columns: tasks). Dashed line marks 0.5. Translated=circles; original=squares. Claude-3.5-sonnet translated runs excluded for xenophobia due to inconsistencies. See Appendix D3 for class-1 F1.

Across both tasks, German yields higher macro-F1 than Danish/Swedish, irrespective of family. The gap is largest on anti-government; on xenophobia it narrows but remains. Danish/Swedish show similar profiles, with xenophobia beats anti-government for every family and method. Fine-tuned encoders hit the highest peaks (often 0.70–0.90+), decoders cluster mid-band (proprietary ≈0.40–0.50+, open-source ≈0.20–0.40), and encoder–decoders are intermediate and variable. Few-shot beats zero-shot for decoders, but gains are modest; translation yields small, mixed effects and does not alter family ranking.

Two qualifications: (1) On anti-government, non-supervised decoders are generally too low for reliable classification; usable performance requires supervision, with encoders setting the ceiling. (2) Danish xenophobia shows several few-shot decoder runs that approach encoders; this does not generalize to German/Swedish.

Taken together, resource level and method interact: German raises the ceiling and amplifies supervision gains; Danish/Swedish benefit most from supervised encoders, while decoder-only models can be pragmatically competitive for Danish xenophobia when fine-tuning is infeasible.

### *Qualitative assessment: Agreement and disagreements between LLMs*

The second part of the analysis examines agreements and disagreements between classifiers to identify linguistic patterns and discourse characteristics that systematically drive divergence between model outputs and expert annotations (RQ2). I coded low-agreement items into six mutually exclusive error types: mediated discourse/reporting, actor-role conflation, numbers-as-argument, irony/sarcasm, euphemism/dog-whistling, and other/ambiguous cases (see Methods and Appendix F for definitions). I compute per-item agreement as the share of valid model predictions matching the gold label and stratify items into low (<.70), mid (.70–<.90),

and high (≥.90) bands (items with <3 valid predictions excluded). See Appendix E-3 for banding rationale and robustness checks.

Table 4 shows a clear task asymmetry in this corpus. Xenophobia exhibits higher consensus (overall high-band 76.6%; 4,598/6,000) and fewer low-agreement items (8.7%; 523/6,000). Anti-government shows a larger low- and mid-band (low 20.3%; 1,216/6,000; mid 27.4%; 1,644/6,000). This should not be read as implying that xenophobia is inherently "overt" and anti-government discourse inherently "covert": both constructs are historically and contextually shaped, and both can be expressed in explicit or coded ways. Rather, in this specific sample, xenophobic posts more often contain clear surface markers (e.g., slurs, exclusionary slogans), whereas anti-government posts more frequently hinge on interpretation of country-specific institutions, policy debates, or ironic framing. This interpretation is consistent with the error analysis, where low-agreement items for both constructs cluster around mediated discourse, numbers-as-argument, irony/sarcasm, and euphemism/dog-whistling, but such cases are relatively more common for the anti-government label.

**Table 4.** Agreement bands across discourse dimensions

| Category | Language | Low Agreement (<70%) | Mid Agreement (70-89%) | High Agreement (≥90%) | Mean Ratio |
|---|---|---|---|---|---|
| **Anti-government** | Danish | 375 (18.8%) | 479 (23.9%) | 1146 (57.3%) | 0.81 |
| **Anti-government** | German | 400 (20.0%) | 556 (27.8%) | 1044 (52.2%) | 0.74 |
| **Anti-government** | Swedish | 441 (22.1%) | 609 (30.4%) | 950 (47.5%) | 0.7 |
| **Anti-government** | All | 1216 (20.3%) | 1644 (27.4%) | 3140 (52.3%) | 0.75 |
| **Xenophobia** | Danish | 166 (8.3%) | 270 (13.5%) | 1564 (78.2%) | 0.89 |
| **Xenophobia** | German | 124 (6.2%) | 249 (12.4%) | 1627 (81.3%) | 0.91 |
| **Xenophobia** | Swedish | 233 (11.7%) | 360 (18.0%) | 1407 (70.3%) | 0.86 |
| **Xenophobia** | All | 523 (8.7%) | 879 (14.6%) | 4598 (76.6%) | 0.89 |

**Note:** N=2,000 posts per language per task. Bands: <.70 (low), .70–<.90 (mid), ≥.90 (high). Mean ratio is the per-item average agreement. English translations were not used to avoid biases in band assignments toward translation quality rather than classifier capability.

Language patterns align with this picture. For xenophobia, German has the highest high-band (81.3%) with Danish close (78.2%); Swedish is lower (70.3%). For anti-government, Swedish records the lowest high-band (47.5%) and highest low-band (22.1%). Language–band distributions differ significantly ($\chi^2$, $p<.001$). A mixed-effects logistic model with crossed item/model intercepts confirms task and language effects: anti-government vs xenophobia OR=0.35 [0.23, 0.52]; German vs Danish OR=0.66 [0.59, 0.74]; Swedish vs Danish OR=0.49 [0.44, 0.55]; all $p<.001$ (see Appendix E4). Item-level SD=2.47 and model-level SD=0.98 indicate substantial text/model heterogeneity; AIC=268,826.

**Error Typology on Low-Agreement Items**

Low-agreement items (< .70; n = 1,727) were coded into six mutually exclusive error types ($\kappa$=0.737): mediated discourse/reporting, actor-role conflation, numbers-as-argument, irony/sarcasm, euphemism/dog-whistling, other/ambiguous (Methods; Appendix F). Figure 2 visualizes distributions by task and language. Tables 5–6 provide examples for anti-government and xenophobia, respectively, highlighting typical triggers (reported speech, sarcasm, euphemism) and brief rationales.

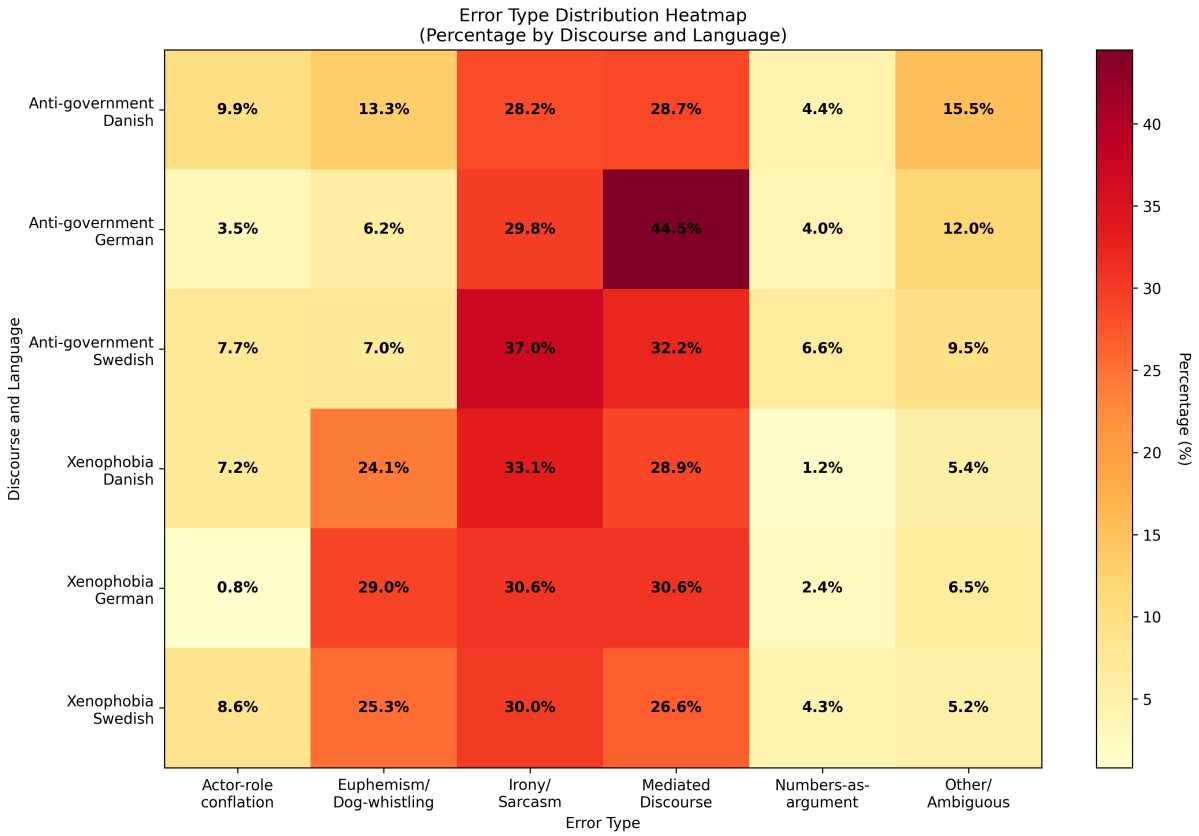


**Figure 2.** Error Type Distribution by Discourse and Language

**Note:** Heatmap showing the percentage distribution of error types across discourse dimensions (anti-government, xenophobia) and languages (Danish, German, Swedish) for low-agreement items (n=1,727). Values represent intercoder agreement on error type classification (κ=0.737).

Anti-government disputes are dominated by mediated discourse/reporting and irony/sarcasm—models misread quoted/attributed speech or neutral headlines as stance and struggle with pragmatic cues (sarcasm, rhetorical address, indirect blame). Xenophobia disagreements center on euphemism/dog-whistling and irony/sarcasm—hostility is communicated via coded language, insinuation, or metaphor, which weakens lexical matching. These qualitative profiles align with the quantitative results (lower agreement odds for anti-government).

**Table 5.** Examples of Disagreement for Anti-government Discourse (<.70)

| Post Text | Language | Error type (agreement ratio) | Rationale |
|---|---|---|---|

| | | | |
|---|---|---|---|
| **The Humorlessness of Annalena Baerbock – Interview with Klemens Kilic. In August 2019, YouTuber Klemens Kilic played a prank on Green Party chair Annalena Baerbock in a selfie video. Baerbock reacted extremely humorlessly: [Link].** | German | Mediated discourse (0.18) | The post mostly reports an interview and links the clip (metadata-driven), with minimal explicit stance; LLMs. may over-index on the word "humorlessness", while the rest reads like straight reporting of a past incident (MD) |
| **"SPD finance minister in Mecklenburg-Western Pomerania resigns."** | German | Mediated discourse (0.19) | Reports on a minister's resignation, without explicit evaluation or another specific error category; little context to anchor tone |
| **"You politicians—and those of you who know nothing, and probably many Muslims too—about what is expected of Danish society when it comes to Islam in public: do yourselves a favor and get some insight 👇"** | Danish | Irony, sarcasm or rhetorical questions (0.31) | Ambiguous target: Uses suggestion ("get some insight") and expectation talk rather than direct claims; models struggle to map insinuation to a category (implied stance) |
| **As usual, it's the naive and easily duped local politicians (probably completely without the proper education!) who have paved the way for this to happen. 'Politicians' who, of course, bear no responsibility!** | Swedish | Irony, sarcasm or rhetorical questions (0.19) | Sarcastic scare quotes around "Politiker."; Classifier uncertainty around local vs. national politicians as anti-government |
| **It has become normal to use the police and the legal system against one's political opponents—both pro-refugee activists and Muslims use this tactic. Here's a citizen who wants to report Rasmus Paludan to the police because his camera brushed past the woman's mailbox.** | Danish | Actor-role conflation (0.24) | The post criticizes pro-refugee activists and Muslims for using the legal system against political opponents, which is a mixing of actor roles that can lead to false labelling. |

Within anti-government, German concentrates errors in mediated discourse/reporting (quotations, attributions), while Swedish shows the highest share of irony/sarcasm, pointing to language specific pragmatics as the main obstacle. Danish splits between the two, indicating both reporting frames and indirect tone contribute substantially. For xenophobia, Danish and Swedish disagreements are chiefly euphemism/dog whistling and irony/sarcasm, whereas German mixes mediated discourse with euphemism/irony, suggesting that both reporting frames and indirectness impede recognition of xenophobic stance.

**Table 6.** Examples of Disagreement for Xenophobic discourse (<.70)

| Post Text | Language | Error type (agreement ratio) | Rationale |
|---|---|---|---|
| **The Alternative's [left-wing party] campaign: Over 1,000 Muslims at a secret Islamic voter meeting. #dkpol #anewwe** | Danish | Euphemism / Dog-whistle (0.19) | whether "secret" is factual or rhetorical; party name "Alternativet" vs generic "alternative"; hashtagged framing. |
| **Listen to this spectacle, I have absolutely no respect for a sect that also does not respect women and animals; Islam has no place whatsoever in a civilized society. So, Margot, with your speech you are nothing but a traitor to the Swedish Christian society.** | Swedish | Euphemism / Dog-whistle (0.31) | Strong invective and religious framing ("sect," "traitor") require tone-sensitive rendering; double commas/ellipses convey rant-like cadence |
| **@FEA66 unfortunately, some Jamaicans have converted to Islam.** | Danish | Euphemism / Dog-whistle (0.23) | understated wording with evaluative "unfortunately" implies stance; very short context makes nuance easy to over/under-interpret. |
| **Where a vacuum arises, Islam pushes in.** | Danish | Irony, sarcasm or rhetorical questions (0.31) | Ambiguous target: Uses suggestion ("get some insight") and expectation talk rather than direct claims; models struggle to map insinuation to a category (implied stance) |
| **As usual, it's the naive and easily duped local politicians (probably completely without the proper education!) who have paved the way for this to happen. 'Politicians' who, of course, bear no responsibility!** | Swedish | Irony, sarcasm or rhetorical questions (0.19) | Sarcastic scare quotes around "Politiker."; Classifier uncertainty around local vs. national politicians as anti-government |

Actor role conflation (who is the government/target) and numbers as argument (statistics used as stance without explicit evaluation) are smaller but consistent contributors – more salient for anti-government (role boundaries) and present for xenophobia when quantitative framing substitutes for overt stance. Overall, disagreements arise when signals are indirect (irony, euphemism), structurally mediated (reporting frames), or require precise role/context disambiguation—conditions under which lexical matching is insufficient and discourse level inference is required.

## Discussion and Limitations

This study asked whether generative LLMs narrow the language gap in multilingual, discourse-level classification and how disagreement with expert labels arises (Baden et al., 2022). Across

33 configurations and two tasks, the answer to RQ1 is clear: resource level dominates. High-resource German (and English translations) consistently raises performance ceilings across families, while low-resource Danish/Swedish lag (macro-F1 encoders ≈0.85 [0.70–0.96] vs decoders ≈0.69–0.70 proprietary; ≈0.63–0.64 open-source). Fine-tuned encoders (mDeBERTa/mmBERT) are strongest overall, especially on the pragmatics-heavy anti-government task, whereas decoder-only models – despite modest few-shot gains – cluster in a lower band and rarely approach supervised encoders. Translation provides small, mixed benefits for prompted decoders and does not change the family ranking, nor does it close the supervision gap.

Turning to RQ2, the consensus audit pinpoints *why* models and experts diverge. Mediated discourse/reporting and irony/sarcasm dominate low-agreement anti-government items; euphemism/dog-whistling plus irony/sarcasm dominate xenophobia disputes. These are exactly the cases where surface lexical cues are weak, stance is attributed rather than asserted, or targets are indirect. The mixed-effects results dovetail: odds of correct classification are markedly lower for anti-government and lower in Swedish than Danish/German. In short, discourse-level inference—not sheer scale—limits reliability. Treating model configurations as raters and banding items by agreement (<.70, .70–<.90, ≥.90) yields a stable sampling frame for error analysis and clarifies where reliability breaks down across languages. Robustness checks (0.20/0.80 and 30–40–30 quantiles) preserve task/language contrasts, but 0.70/0.90 bands maximize interpretability and cell sizes for diagnostics and coding. Class-1 F1 in Appendix D3 further strengthens the encoder advantage, especially on anti-government in German.

Methodologically, treating configurations as raters and banding items by consensus (<.70, .70–<.90, ≥.90; excluding items with <3 valid predictions) provides a principled bridge

between quantitative reliability and qualitative error analysis. Alternative bandings (0.20/0.80; 30–40–30 quantiles) preserve the task/language contrasts, but 0.70/0.90 maximizes interpretability and yields adequate cell sizes for inference and targeted coding. This perspective does not discard disagreement as "noise"; it leverages low-agreement items to expose the pragmatic and structural conditions under which reliability fails—reported speech, role ambiguity, irony/euphemism—thereby creating a reproducible path for diagnosis and improvement.

In practical terms, the results argue for supervised encoders as the default for multilingual political discourse classification. Decoder-only prompting—despite modest few-shot gains and occasional benefits from translation—remains insufficient for complex content such as anti-government discourse. To operationalize reliability at scale, it follows from this study that it is important to organize model development and review around consensus bands: items with low agreement (<.70) are routed to adjudication and error-aware training, mid-band items (.70–<.90) receive targeted quality assurance, and high-band items (≥.90) undergo light audit. Robustness efforts should concentrate on discourse-aware modules that precede classification—detecting reported speech to avoid mistaking quotations for stance, resolving actor–role references to reduce target ambiguity, and handling pragmatic cues such as irony, sarcasm, and euphemism—followed by fine-tuning on error-enriched batches drawn specifically from the low-agreement segment. This workflow links quantitative reliability to concrete model improvements and, in my experiments, yields the clearest path toward raising performance in smaller languages and on the most challenging discourse phenomena.

The paper is not without limitations. *First*, some architecture × language cells include few models (n=1–2), widening uncertainty and limiting in-family rank ordering. Prompting/hyperparameter choices for decoders were standardized to isolate family-level

patterns, but different prompts, sampling, or instruction-tuning could shift absolute scores without changing the observed hierarchy. *Second*, agreement bands (<.70, .70–<.90, ≥.90) are justified empirically and via sensitivity/quantile checks, but any fixed thresholds are conventional. Items with fewer than three valid predictions are excluded; residual missingness could modestly affect band shares. *Third*, findings pertain to two discourse tasks (xenophobia, anti-government) and three languages (Danish, German, Swedish). Generalization beyond these settings should be cautious. Although I treated the human label as gold, the qualitative audit shows that "valid disagreement" can arise under pragmatic indeterminacy (e.g., reported speech, irony, euphemism); consensus bands partly address this, but cannot eliminate subjectivity in all edge cases.

## Conclusion

Multilingual, discourse-level classification requires supervision, resource-sensitive modeling, and discourse awareness. Fine-tuned encoders set the performance ceiling across languages; translation and prompting do not substitute for supervision. Disagreement with expert labels arises chiefly when stance is indirect (irony/euphemism), structurally mediated (reported speech), or role-ambiguous – conditions that surface cues underspecify. By leveraging consensus bands and crossed random-effects modeling, I quantify where and why reliability fails and provide a reproducible path for triage and improvement. The practical recipe is clear: prioritize supervised encoders, organize review and training around consensus bands, and target discourse-level failure modes to raise multilingual reliability in applied political text analysis.

## Data Availability Statement

The annotated dataset is available at ANONYMIZED. Code and materials will be deposited on OSF upon acceptance. Anonymized appendix available at: www.bit.ly/appendix_llm_eval.

## References


Artetxe, M., Goswami, V., Bhosale, S., Fan, A., & Zettlemoyer, L. (2023). Revisiting Machine Translation for Cross-lingual Classification. In H. Bouamor, J. Pino, & K. Bali (Eds.), *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (pp. 6489–6499). Association for Computational Linguistics. https://doi.org/10.18653/v1/2023.emnlp-main.399

AUTHORS. (2024).

AUTHORS. (2025).

Baden, C., Pipal, C., Schoonvelde, M., & van der Velden, M. A. C. G. (2022). Three Gaps in Computational Text Analysis Methods for Social Sciences: A Research Agenda. *Communication Methods and Measures*, *16*(1), 1–18. https://doi.org/10.1080/19312458.2021.2015574

Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D. M., Wu, J., Winter, C., … Amodei, D. (2020). Language Models are Few-Shot Learners. *Advances in Neural Information Processing Systems*, *33*, 1877–1901.

Bucher, M. J. J., & Martini, M. (2024). *Fine-Tuned "Small" LLMs (Still) Significantly Outperform Zero-Shot Generative AI Models in Text Classification* (No. arXiv:2406.08660). arXiv. https://doi.org/10.48550/arXiv.2406.08660

Campion, K., Nolan, M., & O'Brien, N. (2023). Framing the Australian extreme right: Proposing a threefold typology with consideration of legislation and listing regulations. *Journal of Policing, Intelligence and Counter Terrorism*, *18*(3), 282–306. https://doi.org/10.1080/18335330.2023.2165136

Choenni, R., Garrette, D., & Shutova, E. (2023). Cross-Lingual Transfer with Language-Specific Subnetworks for Low-Resource Dependency Parsing. *Computational Linguistics*, *49*(3), 613–641. https://doi.org/10.1162/coli_a_00482

Choenni, R., Rajaee, S., Monz, C., & Shutova, E. (2024). *On the Evaluation Practices in Multilingual NLP: Can Machine Translation Offer an Alternative to Human Translations?* (No. arXiv:2406.14267). arXiv. https://doi.org/10.48550/arXiv.2406.14267

Cisneros, J. D., & Nakayama, T. K. (2015). New Media, Old Racisms: Twitter, Miss America, and Cultural Logics of Race. *Journal of International and Intercultural Communication*, *8*(2), 108–127. https://doi.org/10.1080/17513057.2015.1025328

Conneau, A., Khandelwal, K., Goyal, N., Chaudhary, V., Wenzek, G., Guzmán, F., Grave, E., Ott, M., Zettlemoyer, L., & Stoyanov, V. (2020). Unsupervised Cross-lingual Representation Learning at Scale. In D. Jurafsky, J. Chai, N. Schluter, & J. Tetreault (Eds.), *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (pp. 8440–8451). Association for Computational Linguistics. https://doi.org/10.18653/v1/2020.acl-main.747

Creswell, J. W., & Plano, C. V. L. (2022). *Designing and conducting mixed methods research* (Third edition.). Sage Publications Inc.

Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding* (No. arXiv:1810.04805). arXiv. https://doi.org/10.48550/arXiv.1810.04805

Dong, Q., Li, L., Dai, D., Zheng, C., Ma, J., Li, R., Xia, H., Xu, J., Wu, Z., Liu, T., Chang, B., Sun, X., Li, L., & Sui, Z. (2024). *A Survey on In-context Learning* (No. arXiv:2301.00234). arXiv. https://doi.org/10.48550/arXiv.2301.00234

Eichin, F., Liu, Y. J., Plank, B., & Hedderich, M. A. (2025). *Probing LLMs for Multilingual Discourse Generalization Through a Unified Label Set* (No. arXiv:2503.10515). arXiv. https://doi.org/10.48550/arXiv.2503.10515

ElSherief, M., Ziems, C., Muchlinski, D., Anupindi, V., Seybolt, J., Choudhury, M., & Yang, D. (2021). Latent Hatred: A Benchmark for Understanding Implicit Hate Speech. *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, 345–363. https://doi.org/10.18653/v1/2021.emnlp-main.29

Farjam, M., Meyer, H., & Lohkamp, M. (2025). A Practical Guide and Case Study on How to Instruct LLMs for Automated Coding During Content Analysis. *Social Science Computer Review*, 08944393251349541. https://doi.org/10.1177/08944393251349541

Flach, P. (2019). Performance Evaluation in Machine Learning: The Good, the Bad, the Ugly, and the Way Forward. *Proceedings of the AAAI Conference on Artificial Intelligence*, *33*(01), 9808–9814. https://doi.org/10.1609/aaai.v33i01.33019808

Gilardi, F., Alizadeh, M., & Kubli, M. (2023). *ChatGPT Outperforms Crowd-Workers for Text-Annotation Tasks* (No. arXiv:2303.15056). arXiv. https://doi.org/10.48550/arXiv.2303.15056

Grimmer, J., & Stewart, B. M. (2013). Text as data: The promise and pitfalls of automatic content analysis methods for political texts. *Political Analysis*, *21*(3), 267–297. https://doi.org/10.1093/pan/mps028

Gu, Y., Dong, L., Wei, F., & Huang, M. (2023). Pre-Training to Learn in Context. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 4849–4870. https://doi.org/10.18653/v1/2023.acl-long.267

Guo, Z., Jin, R., Liu, C., Huang, Y., Shi, D., Supryadi, Yu, L., Liu, Y., Li, J., Xiong, B., & Xiong, D. (2023). *Evaluating Large Language Models: A Comprehensive Survey* (No. arXiv:2310.19736). arXiv. https://doi.org/10.48550/arXiv.2310.19736

Heseltine, M., & Clemm Von Hohenberg, B. (2024). Large language models as a substitute for human experts in annotating political text. *Research & Politics*, *11*(1), 20531680241236239. https://doi.org/10.1177/20531680241236239

Hoyos-Osorio, J., Alvarez-Meza, A., Daza-Santacoloma, G., Orozco-Gutierrez, A., & Castellanos-Dominguez, G. (2021). Relevant information undersampling to support imbalanced data classification. *Neurocomputing*, *436*, 136–146. https://doi.org/10.1016/j.neucom.2021.01.033

Kim, Y., & Lee, H. (2025). *Selective Demonstration Retrieval for Improved Implicit Hate Speech Detection* (No. arXiv:2504.12082). arXiv. https://doi.org/10.48550/arXiv.2504.12082

Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., & Iwasawa, Y. (2022). Large Language Models are Zero-Shot Reasoners. *36th Conference on Neural Information Processing Systems*.

Krippendorff, K. (2004). Reliability in Content Analysis.: Some Common Misconceptions and Recommendations. *Human Communication Research*, *30*(3), 411–433. https://doi.org/10.1111/j.1468-2958.2004.tb00738.x

Kristensen-McLachlan, R. D., Canavan, M., Kárdo, M., Jacobsen, M., & Aarøe, L. (2025). Are chatbots reliable text annotators? Sometimes. *PNAS Nexus*, *4*(pgaf069).

Kroon, A., Welbers, K., Trilling, D., & van Atteveldt, W. (2024). Advancing Automated Content Analysis for a New Era of Media Effects Research: The Key Role of Transfer Learning. *Communication Methods and Measures*, *18*(2), 142–162. https://doi.org/10.1080/19312458.2023.2261372

Laurer, M., van Atteveldt, W., Casas, A., & Welbers, K. (2022). *Less Annotating, More Classifying – Addressing the Data Scarcity Issue of Supervised Machine Learning with Deep Transfer Learning and BERT-NLI*.

Laurer, M., van Atteveldt, W., Casas, A., & Welbers, K. (2023). *Lowering the Language Barrier—Investigating Deep Transfer Learning and Machine Translation for Multilingual Analyses of Political Texts*. *5*(2).

Laurer, M., van Atteveldt ,Wouter, Casas ,Andreu, & and Welbers, K. (2025). On Measurement Validity and Language Models: Increasing Validity and Decreasing Bias with Instructions. *Communication Methods and Measures*, *19*(1), 46–62. https://doi.org/10.1080/19312458.2024.2378690

Lauscher, A., Ravishankar, V., Vulić, I., & Glavaš, G. (2020). *From Zero to Hero: On the Limitations of Zero-Shot Cross-Lingual Transfer with Multilingual Transformers* (No. arXiv:2005.00633). arXiv. https://doi.org/10.48550/arXiv.2005.00633

Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., Stoyanov, V., & Zettlemoyer, L. (2020). BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension. In D. Jurafsky, J. Chai, N. Schluter, & J. Tetreault (Eds.), *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (pp. 7871–7880). Association for Computational Linguistics. https://doi.org/10.18653/v1/2020.acl-main.703

Licht, H. (2023). Cross-Lingual Classification of Political Texts Using Multilingual Sentence Embeddings. *Political Analysis*, 1–14. https://doi.org/10.1017/pan.2022.29

Licht, H., Sczepanski, R., Laurer, M., & Bekmuratovna, A. (2024). *No more cost in translation: Validating open-source machine translation for quantitative text analysis*. Open Science Framework. https://doi.org/10.31219/osf.io/9trjs

Matamoros-Fernández, A., & Farkas, J. (2021). Racism, Hate Speech, and Social Media: A Systematic Review and Critique. *Television & New Media*, *22*(2), 205–224. https://doi.org/10.1177/1527476420982230

Máté, Á., Sebők, M., Wordliczek, Ł., Stolicki, D., & Feldmann, Á. (2023). Machine Translation as an Underrated Ingredient? Solving Classification Tasks with Large Language Models for Comparative Research. *Computational Communication Research*, *5*(2). https://doi.org/HTTPS://DOI.ORG/10.5117/CCR2023.2.6.MATE

Mets, M., Karjus, A., Ibrus, I., & Schich, M. (2024). Automated stance detection in complex topics and small languages: The challenging case of immigration in polarizing news media. *PLOS ONE*, *19*(4), e0302380. https://doi.org/10.1371/journal.pone.0302380

Mudde, C. (1995). Right-wing extremism analyzed: A comparative analysis of the ideologies of three alleged right-wing extremist parties (NPD, NDP, CP'86). *European Journal of Political Research*, *27*(2), 203–224. https://doi.org/10.1111/j.1475-6765.1995.tb00636.x

Nassiri, K., & Akhloufi, M. (2023). Transformer models used for text-based question answering systems. *Applied Intelligence*, *53*(9), 10602–10635. https://doi.org/10.1007/s10489-022-04052-8

Nielsen, D. S., Enevoldsen, K., & Schneider-Kamp, P. (2025). *Encoder vs Decoder: Comparative Analysis of Encoder and Decoder Language Models on Multilingual NLU Tasks* (No. arXiv:2406.13469). arXiv. https://doi.org/10.48550/arXiv.2406.13469

Ocampo, N., Sviridova, E., Cabrio, E., & Villata, S. (2023). An In-depth Analysis of Implicit and Subtle Hate Speech Messages. *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, 1997–2013. https://doi.org/10.18653/v1/2023.eacl-main.147

Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Askell, A., Welinder, P., Christiano, P., Leike, J., & Lowe, R. (2022). *Training language models to follow instructions with human feedback* (No. arXiv:2203.02155). arXiv. https://doi.org/10.48550/arXiv.2203.02155

Pires, T., Schlinger, E., & Garrette, D. (2019). *How multilingual is Multilingual BERT?* (No. arXiv:1906.01502). arXiv. https://doi.org/10.48550/arXiv.1906.01502

Radford, A., Narasimhan, K., Salimans, T., & Sutskever, I. (2018). *Improving Language Understanding by Generative Pre-Training*.

Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., & Liu, P. J. (2020). *Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer*.

Rathje, S., Mirea, D.-M., Sucholutsky, I., Marjieh, R., Robertson, C. E., & Van Bavel, J. J. (2024). GPT is an effective tool for multilingual psychological text analysis. *Proceedings of the National Academy of Sciences*, *121*(34), e2308950121. https://doi.org/10.1073/pnas.2308950121

Renze, M., & Guven, E. (2024). The Effect of Sampling Temperature on Problem Solving in Large Language Models. *Findings of the Association for Computational Linguistics: EMNLP 2024*, 7346–7356. https://doi.org/10.18653/v1/2024.findings-emnlp.432

Robinson, N. R., Ogayo, P., Mortensen, D. R., & Neubig, G. (2023). *ChatGPT MT: Competitive for High- (but not Low-) Resource Languages* (No. arXiv:2309.07423). arXiv. https://doi.org/10.48550/arXiv.2309.07423

Rousseeuw, P. J., & Leroy, A. M. (1987). *Robust Regression and Outlier Detection* (1st ed.). Wiley. https://doi.org/10.1002/0471725382

Schick, T., & Schütze, H. (2021). *Exploiting Cloze Questions for Few Shot Text Classification and Natural Language Inference* (No. arXiv:2001.07676). arXiv. https://doi.org/10.48550/arXiv.2001.07676

Sivarajkumar, S., Kelley, M., Samolyk-Mazzanti, A., Visweswaran, S., & Wang, Y. (2024). An Empirical Evaluation of Prompting Strategies for Large Language Models in Zero-Shot Clinical Natural Language Processing: Algorithm Development and Validation Study. *JMIR Medical Informatics*, *12*, e55318. https://doi.org/10.2196/55318

Sun, C., Qiu, X., Xu, Y., & Huang, X. (2020). *How to Fine-Tune BERT for Text Classification?* (No. arXiv:1905.05583). arXiv. http://arxiv.org/abs/1905.05583

Törnberg, P. (2024). Large Language Models Outperform Expert Coders and Supervised Classifiers at Annotating Political Social Media Messages. *Social Science Computer Review*, 08944393241286471. https://doi.org/10.1177/08944393241286471

Weber, M., & Reichardt, M. (2023). *Evaluation is all you need. Prompting Generative Large Language Models for Annotation Tasks in the Social Sciences. A Primer using Open Models* (No. arXiv:2401.00284). arXiv. https://doi.org/10.48550/arXiv.2401.00284

Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q., & Zhou, D. (2023). *Chain-of-Thought Prompting Elicits Reasoning in Large Language Models* (No. arXiv:2201.11903). arXiv. https://doi.org/10.48550/arXiv.2201.11903

Wei, L., Li, L., Xiang, T., Liu, X., & Garcia, N. (2025). Cracking the Code: Enhancing Implicit Hate Speech Detection through Coding Classification. *ArXiv*.

Windsor, L. C., Cupit, J. G., & Windsor, A. J. (2019). Automated content analysis across six languages. *PLOS ONE*, *14*(11), e0224425. https://doi.org/10.1371/journal.pone.0224425

Wu, S., & Dredze, M. (2020). *Are All Languages Created Equal in Multilingual BERT?* (No. arXiv:2005.09093). arXiv. http://arxiv.org/abs/2005.09093

Xue, L., Constant, N., Roberts, A., Kale, M., Al-Rfou, R., Siddhant, A., Barua, A., & Raffel, C. (2021). *mT5: A massively multilingual pre-trained text-to-text transformer* (No. arXiv:2010.11934). arXiv. http://arxiv.org/abs/2010.11934

Zhao, W. X., Zhou, K., Li, J., Tang, T., Wang, X., Hou, Y., Min, Y., Zhang, B., Zhang, J., Dong, Z., Du, Y., Yang, C., Chen, Y., Chen, Z., Jiang, J., Ren, R., Li, Y., Tang, X., Liu, Z., … Wen, J.-R. (2023). *A Survey of Large Language Models* (Version 16). arXiv. https://doi.org/10.48550/ARXIV.2303.18223

Zhu, X., Gardiner, S., Roldán, T., & Rossouw, D. (2024). The Model Arena for Cross-lingual Sentiment Analysis: A Comparative Study in the Era of Large Language Models. *Proceedings of the 14th Workshop on Computational Approaches to Subjectivity, Sentiment, & Social Media Analysis*, 141–152. https://doi.org/10.18653/v1/2024.wassa-1.12

Zhu, Y., Moniz, J. R. A., Bhargava, S., Lu, J., Piraviperumal, D., Li, S., Zhang, Y., Yu, H., & Tseng, B.-H. (2024). *Can Large Language Models Understand Context?* (No. arXiv:2402.00858). arXiv. https://doi.org/10.48550/arXiv.2402.00858